

Session 5: Building AI-Powered Java Apps

Spring AI, RAG, Tools, and MCP

Press Space for next page →



Welcome to Session 5!

Building AI-Powered Java Applications

From Spring Boot to Spring AI

Spring Boot 3.5 + Spring AI 1.1.0 + Java 21

What We'll Build Today

- **Spring AI ChatClient** - Fluent API for LLM interactions
- **Prompt Templates** - Reusable, parameterized prompts
- **RAG Pipeline** - Chat with your documents
- **Function Calling** - Give AI tools to execute code
- **MCP Integration** - Enhanced context for Cursor

Course Journey

- **Session 1:** Cursor fundamentals
- **Session 2:** Mobile development with AI
- **Session 3:** Agentic coding patterns
- **Session 4:** AI-assisted testing
- **Session 5:** Building AI apps with Spring AI ← Today

Today's Stack

Spring Boot 3.5.7 • Spring AI 1.1.0 • Java 21

- OpenAI or Anthropic API keys required
- All code available in `spring-ai-demo/` folder
- Labs guide you through each feature

Part 1: Introduction to Spring AI

The Spring Way to Build AI Applications

- Official Spring project for AI integration
- Portable abstraction over AI providers
- Familiar Spring programming model



What is Spring AI?

- **Official Spring Project** for AI integration
- **Portable abstraction** over AI providers
- **Spring Boot auto-configuration**
- **Familiar Spring programming model**

Spring AI Core Components

- **ChatClient:** Fluent API for LLM interactions
- **Embeddings:** Vector representations of text
- **Vector Stores:** Storage for document embeddings
- **Function Calling:** Tools that AI can invoke
- **Document Readers:** PDF, Word, text processing

Spring AI Advantages

- Switch between OpenAI, Anthropic, Ollama without code changes
- Dependency injection for AI components
- Spring Boot conventions and auto-configuration
- Familiar patterns: RestTemplate → ChatClient

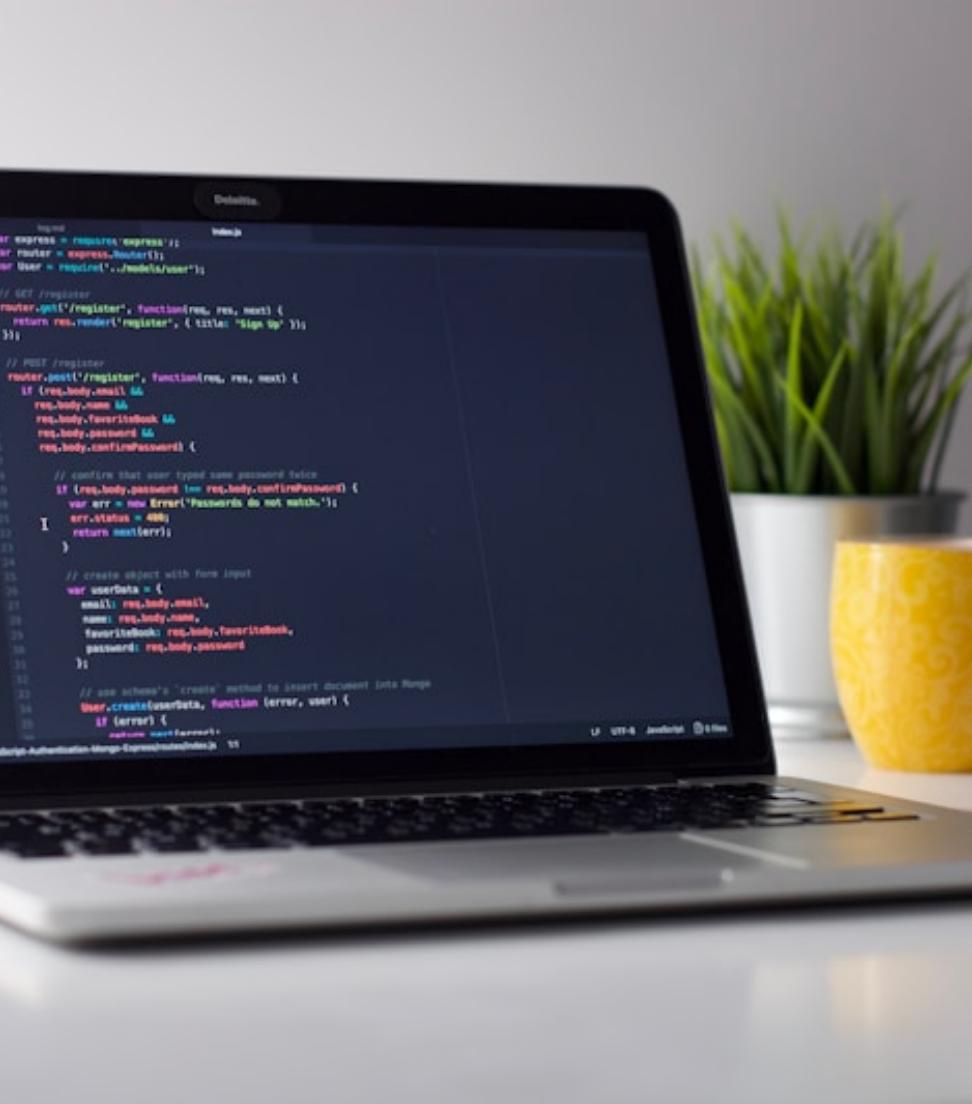
Spring AI Maven Dependencies

```
1 <parent>
2   <groupId>org.springframework.boot</groupId>
3   <artifactId>spring-boot-starter-parent</artifactId>
4   <version>3.5.7</version>
5 </parent>
6
7 <dependencies>
8   <dependency>
9     <groupId>org.springframework.ai</groupId>
10    <artifactId>spring-ai-starter-model-openai</artifactId>
11  </dependency>
12 </dependencies>
```

Spring AI Configuration

```
1 # application.properties
2 spring.ai.openai.api-key=${OPENAI_API_KEY}
3 spring.ai.openai.chat.options.model=gpt-4o
4 spring.ai.openai.chat.options.temperature=0.7
```

- Set `OPENAI_API_KEY` environment variable
- Or use `.env` file with Spring Boot
- Alternative: Use Anthropic with `spring-ai-anthropic`



Part 2: Chat Client & Templating

Fluent API for LLM Interactions

- Build requests with fluent builder pattern
- System and user message configuration
- Prompt templates for reusable prompts

ChatClient Basics

```
1  @RestController
2  @RequestMapping("/api/chat")
3  public class ChatController {
4
5      private final ChatClient chatClient;
6
7      public ChatController(ChatClient.Builder builder) {
8          this.chatClient = builder.build();
9      }
10
11     @GetMapping
12     public String chat(@RequestParam String message) {
13         return chatClient.prompt()
14             .user(message)
15             .call()
16             .content();
17     }
18 }
```

ChatClient Features

- **Fluent API** for building requests
- **System and user messages** configuration
- **Response parsing** and handling
- **Streaming responses** (optional)

Agent Mode Prompt:

```
1 Create a ChatController with a GET endpoint /chat.  
2 Inject ChatClient.Builder and return LLM response.
```

System Prompts

```
1  @GetMapping("/expert")
2  public String expertChat(@RequestParam String topic) {
3      return chatClient.prompt()
4          .system("""
5              You are an expert software architect
6              specializing in Spring Boot applications.
7              Provide concise, practical advice.
8              """)
9      .user("How do I implement " + topic)
10     .call()
11     .content();
12 }
```

Structured Responses

```
1 record BookReview(String title, int rating, String summary) {}
2
3 @GetMapping("/review")
4 public BookReview getBookReview(@RequestParam String book) {
5     return chatClient.prompt()
6         .user("Write a review of the book: " + book)
7         .call()
8         .entity(BookReview.class);
9 }
```

Spring AI automatically: Generates JSON schema → Instructs LLM → Parses to Java object

Prompt Templates

Template File: src/main/resources/prompts/joke.st

- ```
1 Tell me a {style} joke about {topic}.
2 Make it appropriate for a professional audience.
```

## Key Points:

- StringTemplate format (.st files)
- Variable substitution with {variableName}
- Version control your prompts

# Using Prompt Templates

```
1 @GetMapping("/joke")
2 public String tellJoke(
3 @RequestParam String topic,
4 @RequestParam(defaultValue = "funny") String style) {
5
6 return chatClient.prompt()
7 .user(u -> u.text(
8 "classpath:/prompts/joke.st",
9 Map.of("topic", topic, "style", style)
10))
11 .call()
12 .content();
13 }
```

# Student Exercise: Chat Interface

Time: 10 minutes

1. **Create** a new controller
2. **Inject** ChatClient.Builder
3. **Add** a system prompt for your domain
4. **Create** a template for common queries
5. **Test** with various inputs

# Challenge Ideas

- **Code reviewer** - Analyze Java code snippets
- **Documentation generator** - Create JavaDoc
- **SQL translator** - Natural language to SQL
- **Tech explainer** - Simplify complex topics

**Bonus:** Use `.entity()` to return structured responses as Java records

# Part 3: Retrieval Augmented Generation (RAG)

## Chat with Your Documents

- Ground AI responses in your data
- Vector stores for semantic search
- Document chunking and embeddings



# Why RAG?

- **Ground AI in your data** - Not just training data
- **Prevent hallucinations** - Provide context
- **Domain-specific knowledge** - Your documents, policies, code
- **Up-to-date information** - Add new docs anytime

# RAG Pipeline

# RAG Key Concepts

- **Chunking:** Split documents into manageable pieces
- **Embeddings:** Convert text to vectors (meaning as numbers)
- **Vector Similarity:** Find chunks semantically similar to query
- **Context Injection:** Add retrieved chunks to LLM prompt

# Document Ingestion

```
1 @Component
2 public class DocumentLoader implements CommandLineRunner {
3
4 private final VectorStore vectorStore;
5
6 public DocumentLoader(VectorStore vectorStore) {
7 this.vectorStore = vectorStore;
8 }
9
10 @Override
11 public void run(String... args) {
12 Resource resource = new ClassPathResource("documents/policy.txt");
13 TextSplitter splitter = new TokenTextSplitter();
14 List<Document> documents = splitter.split(new TextReader(resource).get());
```

# Document Ingestion (continued)

```
1 // Generate embeddings and store
2 vectorStore.add(documents);
3 log.info("Loaded {} documents", documents.size());
4 }
5 }
```

## What Happens:

1. Load `policy.txt` from classpath
2. Split into ~500 token chunks
3. Generate embeddings (via OpenAI)
4. Store vectors in SimpleVectorStore

# SimpleVectorStore (In-Memory)

```
1 @Configuration
2 public class VectorStoreConfig {
3
4 @Bean
5 public VectorStore vectorStore(EmbeddingModel embeddingModel) {
6 return new SimpleVectorStore(embeddingModel);
7 }
8 }
```

**Good for:** Development, testing, small document sets, prototypes

# Production Vector Stores

- **Chroma** - Open source, easy setup
- **Pinecone** - Managed service
- **PgVector** - PostgreSQL extension
- **Redis** - If already using Redis

```
1 @Bean
2 public VectorStore vectorStore(JdbcTemplate jdbc, EmbeddingModel model) {
3 return new PgVectorStore(jdbc, model);
4 }
```

# RAG Controller Setup

```
1 @RestController
2 @RequestMapping("/api/rag")
3 public class RagController {
4
5 private final ChatClient chatClient;
6 private final VectorStore vectorStore;
7
8 public RagController(ChatClient.Builder builder, VectorStore vectorStore) {
9 this.chatClient = builder.build();
10 this.vectorStore = vectorStore;
11 }
```

# RAG Query Implementation

```
1 @GetMapping("/query")
2 public String query(@RequestParam String question) {
3 // Search for similar documents
4 List<Document> similarDocs = vectorStore.similaritySearch(
5 SearchRequest.query(question).withTopK(5)
6);
7
8 String context = similarDocs.stream()
9 .map(Document::getContent)
10 .collect(Collectors.joining("\n\n"));
```

# RAG Response Generation

```
1 return chatClient.prompt()
2 .system("""
3 Answer the question based ONLY on the provided context.
4 If you cannot answer from the context, say so.
5 Context: {context}
6 """)
7 .user(question)
8 .call()
9 .content();
10 }
11 }
```

# Chunking Strategies

- **TokenTextSplitter** - By token count (most common)
- **Paragraph splitter** - Natural boundaries
- **Sliding window** - Overlap for continuity

**Sweet spot:** 300-800 tokens with 10-20% overlap

# Search Configuration

```
1 SearchRequest.query(question)
2 .withTopK(5) // Return top 5 matches
3 .withSimilarityThreshold(0.7) // Min similarity
4 .withFilterExpression("type == 'policy'"); // Metadata filter
```

- Explicitly tell AI to use context
- Handle "I don't know" gracefully
- Test with questions not in docs

# Student Exercise: RAG Pipeline

Time: 20 minutes

1. **Create** sample documents in `src/main/resources/documents/`
2. **Configure** VectorStore bean
3. **Implement** DocumentLoader
4. **Create** RAG endpoint
5. **Test** with queries requiring document knowledge

# RAG Challenge Ideas

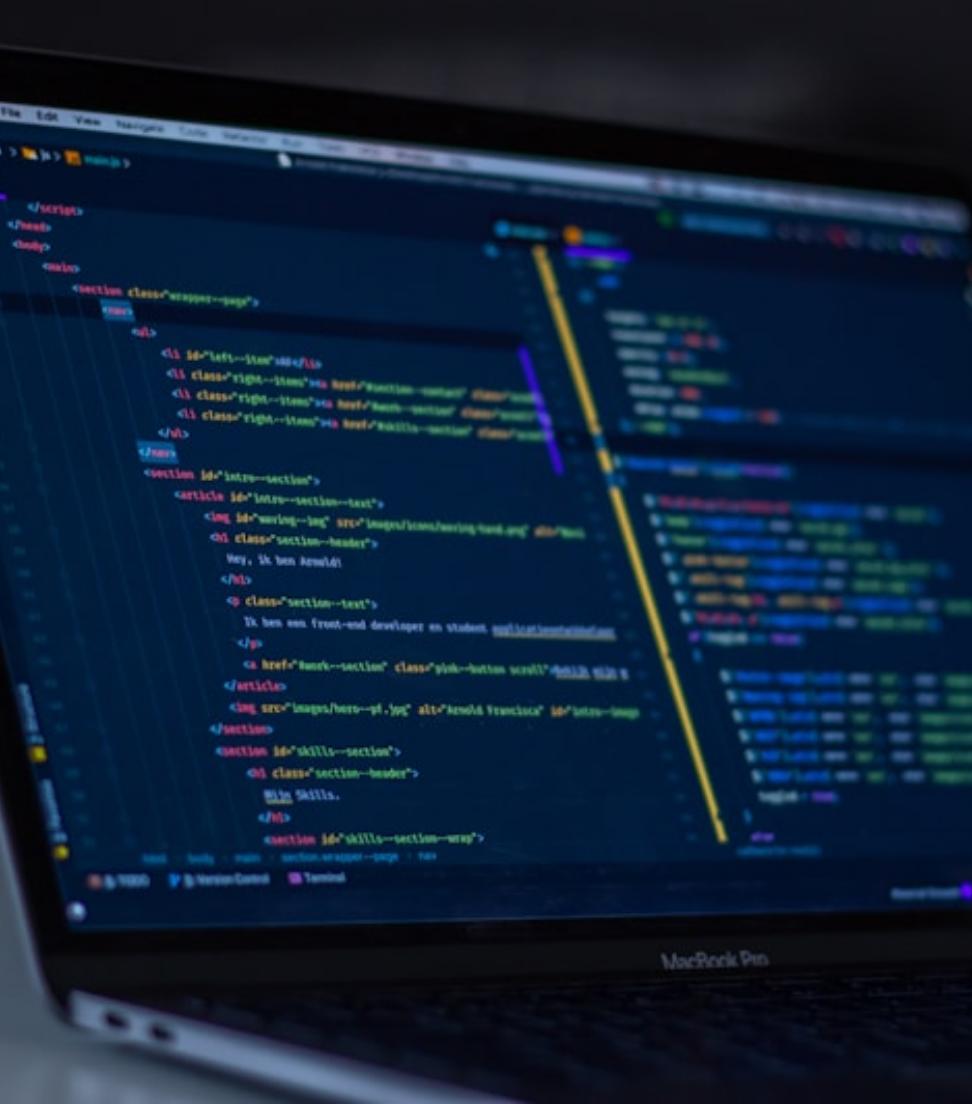
- **Company policies** - HR handbook, procedures
- **Technical docs** - API documentation
- **Knowledge base** - FAQ, troubleshooting

**Bonus:** Add metadata (author, date) and filter search results

# BREAK

10 minutes

Grab coffee before we dive into function calling!



# Part 4: Tools & Function Calling

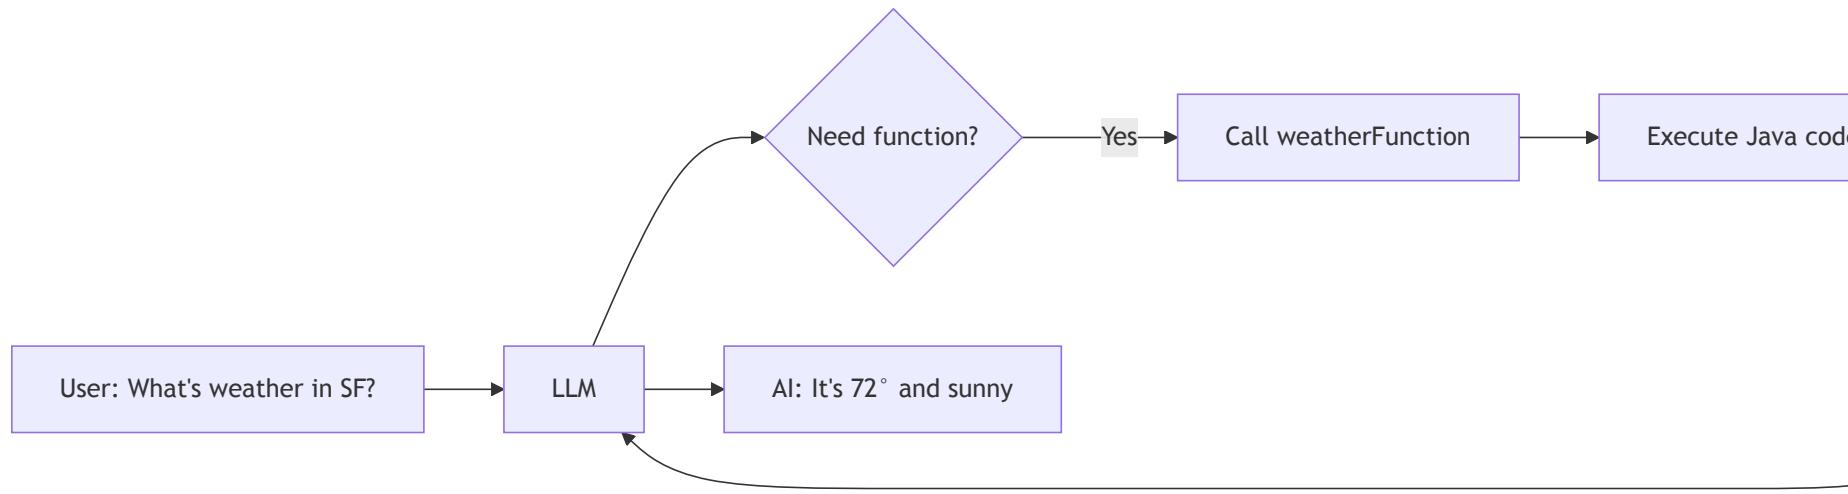
Give AI the Ability to Execute Code

- AI decides when to call functions
- Structured output for function invocation
- Multi-step workflow automation

# What is Function Calling?

- **AI decides** when to call functions
- **Structured output** from LLM (function name + args)
- **Your code executes** the function
- **Return result** to AI for response

# Function Calling Flow



# Function Calling Use Cases

- **Database lookups** - Query data
- **API calls** - External services
- **Calculations** - Math, business logic
- **Workflow automation** - Multi-step tasks

**Key:** AI can call your Java methods based on natural language!

# Weather Tool: Request/Response

```
1 record WeatherRequest(
2 @JsonProperty(required = true, value = "location") String location,
3 @JsonProperty(required = false, value = "unit") String unit
4) {}
5
6 record WeatherResponse(
7 String location, String temperature, String description
8) {}
```

# Weather Tool: Function Bean

```
1 @Configuration
2 public class ToolConfig {
3
4 @Bean
5 @Description("Get current weather for a location")
6 public Function<WeatherRequest, WeatherResponse> weatherFunction() {
7 return request -> new WeatherResponse(
8 request.location(), "72°F", "Sunny with light clouds"
9);
10 }
11 }
```

**Key:** `@Description` tells AI when to call this function

# Register Functions with ChatClient

```
1 @RestController
2 @RequestMapping("/api/tools")
3 public class ToolController {
4
5 private final ChatClient chatClient;
6
7 public ToolController(ChatClient.Builder builder,
8 Function<WeatherRequest, WeatherResponse> weatherFunction) {
9 this.chatClient = builder
10 .defaultFunctions(weatherFunction)
11 .build();
12 }
13
14 @GetMapping("/chat")
15 public String chat(@RequestParam String message) {
16 return chatClient.prompt().user(message).call().content();
17 }
18 }
```

# Function Calling in Action

User: "What's the weather in San Francisco?"

1. AI recognizes need for `weatherFunction`
2. Extracts location: "San Francisco"
3. Calls your Java function
4. Receives: 72°F, Sunny
5. Responds: "It's currently 72°F and sunny..."

# Database Access Tool

```
1 @Entity
2 public class User {
3 @Id private Long id;
4 private String email;
5 private String firstName, lastName;
6 }
7
8 public interface UserRepository extends JpaRepository<User, Long> {
9 Optional<User> findByEmail(String email);
10 }
```

# User Lookup Function Bean

```
1 @Bean
2 @Description("Find user by email address")
3 public Function<UserLookupRequest, UserLookupResponse>
4 userLookupFunction(UserRepository repo) {
5
6 return request -> repo.findByEmail(request.email())
7 .map(u -> new UserLookupResponse(
8 u.getFirstName() + " " + u.getLastName(), u.getEmail()
9))
10 .orElse(new UserLookupResponse("Unknown", request.email()));
11 }
```

# Tool Design Best Practices

- **Clear descriptions** - AI needs to understand purpose
- **Strong typing** - Use records for parameters
- **Validation** - Check inputs before executing
- **Error handling** - Return meaningful errors

# Good Description Examples

```
1 @Description("Get user by email. Returns name and email if found.")
2
3 @Description("Calculate order total including tax. Returns USD.")
4
5 @Description("Send email notification. Returns success status.")
```

# Security Considerations

- **Authentication** - Verify user context
- **Authorization** - Check permissions
- **Input validation** - Sanitize all inputs
- **Audit logging** - Track function calls

```
1 @Bean
2 public Function<OrderLookup, OrderResponse> orderFunction(
3 OrderService service, SecurityContext security) {
4 return request -> {
5 security.checkPermission("orders:read");
6 return service.getOrder(request.orderId());
7 };
8 }
```

# Student Exercise: Custom Tools

Time: 20 minutes

1. **Choose** a domain (e-commerce, HR, etc.)
2. **Create** request/response records
3. **Implement** function bean with @Description
4. **Register** with ChatClient
5. **Test** with natural language queries

# Tool Challenge Ideas

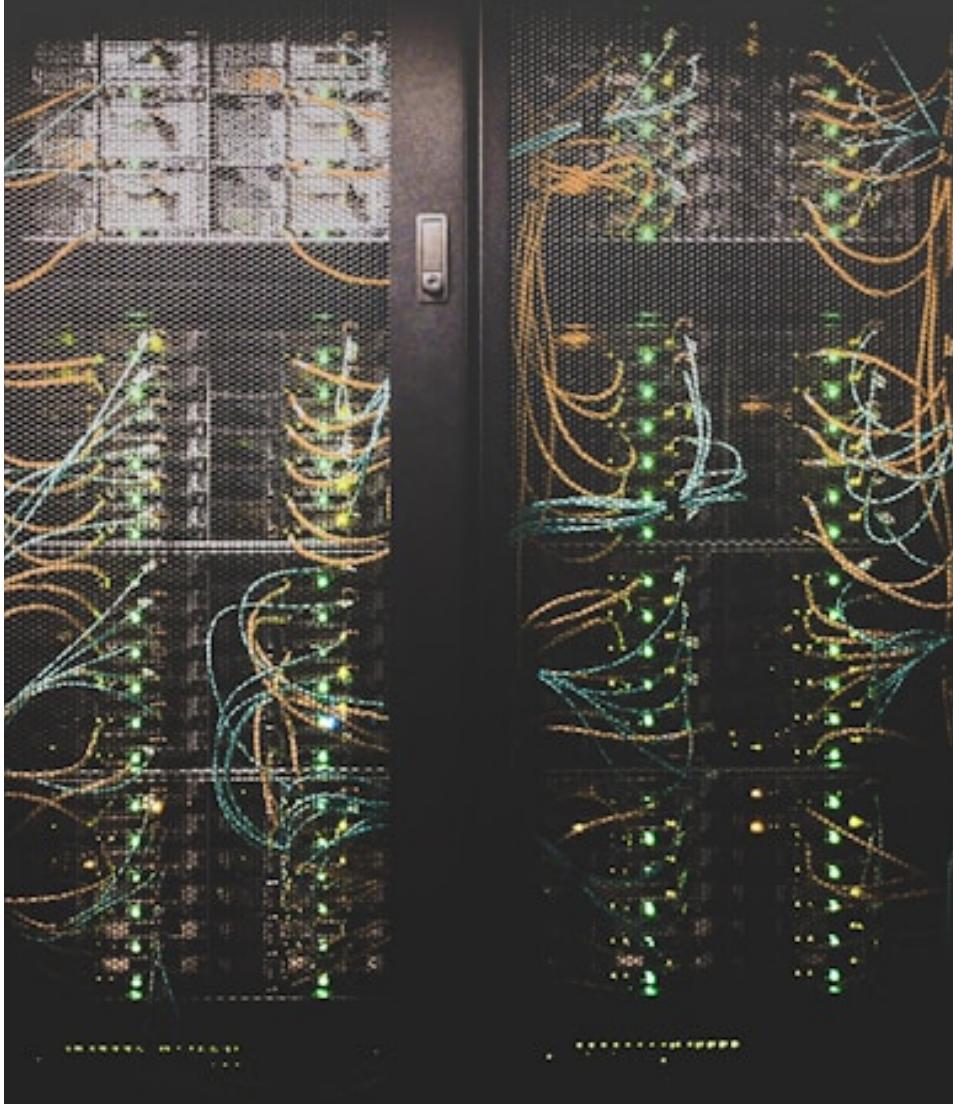
- **Calculator** - Math operations
- **Time converter** - Timezones, formats
- **Currency exchange** - Convert currencies
- **Database query** - Look up records

**Bonus:** Create related functions (createOrder + getOrderStatus + cancelOrder)

# Part 5: Model Context Protocol (MCP)

## Enhanced Context for Cursor

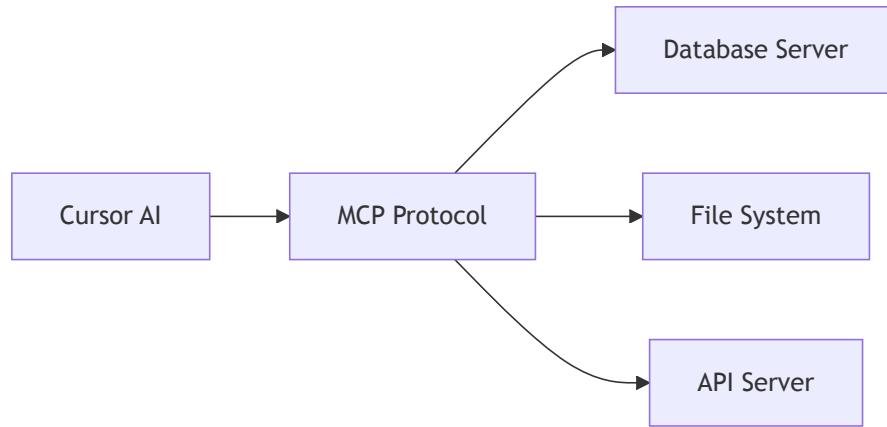
- Protocol for connecting AI to external data
- Real-time database schema awareness
- Tool discovery and dynamic resources



# What is MCP?

- **Protocol** for connecting AI to external data
- **Cursor's MCP support** - Enhanced context in Cursor
- **Standard interface** for tool integration
- **Resources** - Files, databases, APIs

# MCP Architecture



# MCP Benefits

- **Real-time context** - Current database schema
- **Dynamic resources** - Access live data
- **Better suggestions** - More accurate code generation
- **Reduced hallucinations** - Grounded in actual data

# MCP Setup in Cursor

**Step 1:** Cursor → Settings → Features → MCP

```
1 {
2 "mcpServers": {
3 "database": {
4 "command": "npx",
5 "args": ["-y", "@modelcontextprotocol/server-postgres",
6 "postgresql://localhost/mydb"]
7 }
8 }
9 }
```

**Step 2:** Restart Cursor

# Available MCP Servers

## Official:

- `server-postgres` - PostgreSQL
- `server-filesystem` - File access
- `server-github` - GitHub API

## Community:

- Jira, Slack, Notion integrations
- Cloud providers (AWS, GCP, Azure)

# Spring AI + MCP

**Current Approach:** Use Spring AI functions as MCP tool building blocks

```
1 @Bean
2 public Function<SchemaRequest, SchemaResponse>
3 getDatabaseSchema(DataSource dataSource) {
4 return request -> {
5 // Query information_schema
6 // Return table/column details
7 };
8 }
```

# MCP Use Cases

- **Database schema** - Generate accurate SQL, create JPA entities
- **API documentation** - Correct endpoint usage, auth patterns
- **Codebase navigation** - Find related code, understand structure
- **Real-time data** - Status, metrics, analytics

# Student Exercise: MCP Exploration

Time: 10 minutes

1. **Open** Cursor Settings → Features → MCP
2. **Add** an MCP server (filesystem or database)
3. **Test** by asking Cursor about your data
4. **Observe** improved context in responses

# MCP Exploration Questions

- How does MCP affect Cursor's suggestions?
- What queries work better with MCP?
- Can you ask about database structure?
- Does Cursor understand your project better?

**Try:** "What are the main entities in my database?"



# Part 6: Legacy Modernization with AI

## Applying AI Patterns to Legacy Code

- AI-assisted code analysis
- Incremental refactoring strategies
- Adding AI capabilities to existing systems

# Legacy Code Analysis

```
1 // Legacy OrderService (Spring Boot 2.7)
2 @Service
3 public class OrderService {
4 @Autowired private OrderRepository orderRepo;
5 @Autowired private UserRepository userRepo;
6
7 public void processOrder(Long orderId) throws Exception {
8 Order order = orderRepo.findById(orderId).orElseThrow();
9 // Complex business logic, no error handling, hard to test
10 }
11 }
```

# AI Identifies Issues

- Field injection anti-pattern
- Missing validation
- Poor error handling
- No transaction management
- Opportunities for AI enhancement

**Prompt:** "Analyze this legacy OrderService and create a migration plan"

# Step 1: Constructor Injection

```
1 public OrderService(OrderRepository orderRepo, UserRepository userRepo) {
2 this.orderRepo = orderRepo;
3 this.userRepo = userRepo;
4 }
```

## Step 2: Proper Error Handling

```
1 public Order processOrder(Long orderId) {
2 return orderRepo.findById(orderId)
3 .map(this::validateAndProcess)
4 .orElseThrow(() -> new OrderNotFoundException(orderId));
5 }
```

## Step 3: Add AI Capabilities

```
1 public OrderAnalysisReport analyzeOrder(Long orderId) {
2 Order order = getOrder(orderId);
3
4 return chatClient.prompt()
5 .system("Analyze order for risks")
6 .user(toJson(order))
7 .call()
8 .entity(OrderAnalysisReport.class);
9 }
```

# Student Exercise: Legacy Analysis

**Time:** 10 minutes - Use Extended Thinking to explore:

1. "What are the main testing challenges in this legacy codebase?"
2. "How would you prioritize testing improvements?"
3. "What risks should be considered when adding tests?"

# AI-Powered Modernization

- Identify code smells and anti-patterns
- Suggest refactoring strategies
- Generate tests for legacy code
- Add AI capabilities to existing features

**AI suggests:** RAG for policy lookup, function calling for external services

# Wrap-Up & Next Steps

Course Completion

# Spring AI Decision Tree

- **Chat interface?** → Use ChatClient with templates
- **Chat with your data?** → Implement RAG pipeline
- **AI execute code?** → Use function calling
- **Enhanced context?** → Configure MCP

# What We Accomplished

- Spring AI application setup
- ChatClient with prompt templates
- RAG pipeline for document Q&A
- Function calling with Spring AI tools
- MCP exploration for enhanced context

# Core Pattern

Spring AI brings AI capabilities with familiar patterns:

`RestTemplate` → `ChatClient`

`JpaRepository` → `VectorStore`

`@Bean` → `@Bean Function`

# The Five-Session Arc

- **Session 1:** Cursor fundamentals - Chat, Agent, Composer
- **Session 2:** Mobile development - Kotlin, Jetpack Compose
- **Session 3:** Agentic coding patterns
- **Session 4:** AI-assisted testing - JUnit 5, Mockito, TestContainers
- **Session 5:** Spring AI - RAG, Function Calling, MCP

# Next Steps

- **Apply Spring AI** to your projects
- **Experiment** with different AI providers
- **Build** custom tools for your domain
- **Stay updated** on Spring AI releases

# Resources

- [Spring AI Docs](#)
- [Spring AI GitHub](#)
- [Spring AI Examples](#)
- [MCP Specification](#)

# Questions & Discussion

Spring AI • RAG • Function Calling • MCP

# Thank You!

## Building AI-Powered Java Apps with Spring AI

**You're now equipped to build intelligent Java applications!**

Ready for the labs? Let's build with Spring AI!